

Milvus Paper sharing

Product quantization for nearest neighbor search

godchen fishpenguin

Content

1. Background

2. Product quantizers

3. ADC vs SDC

4. IVFADC

5. Experiments

background

K Dimension-tree

KD-tree [5] have been proposed to reduce the search time. However, for high dimensions it turns out [6] that such approaches are not more efficient than the brute-force exhaustive distance calculation, whose complexity is $O(nD)$.

LSH(locality sensitive hashing)

In the case of E2LSH, the memory usage may even be higher than that of the original vectors. Moreover, E2LSH need to perform a final re-ranking step based on exact L2 distances, which requires the indexed vectors to be stored in main memory if access speed is important. It performs well in small and middle size datasets(up to tens of millions)

Product quantizers

i. Vector quantization

Purpose: reduce the cardinality of the representation space, in particular when the input data is real-valued

K-Means

Lloyd optimality conditions:

- (1) a vector x must be quantized to its nearest codebook centroid;
- (2) the reconstruction value must be the expectation of the vectors lying in the Voronoi cell.

ii. Product quantizers

Trade off between search quality and memory requirements.

Product quantizers

数据集大小为N
向量维度为D
 $N * D$




维度为D的向量
划分为m组



对m组向量分别进
行Kmeans聚类



cluster个数 $K*m$
每组有K个D/m维
的聚类中心



每个子向量，用其所属的cluster_id表示
压缩前数据大小为 $N*D*4*8\text{bits (float)}$
压缩后 压缩后 $N * \log_2 K * m$

| | | | | |
|-----|-----|-----|-----|-----|
| 1 | 2 | 2 | ... | 3 |
| 208 | 99 | 91 | ... | 128 |
| 211 | 108 | 79 | ... | 139 |
| ... | ... | ... | ... | ... |
| 109 | 28 | 65 | ... | 89 |



~~$k \times D$~~



$$(k^*)^m = k$$


~~$m \times k^* \times D$~~

~~$* = k^* \times D$~~

~~$= D \times k^{1/m}$~~

m


Product quantizers




ADC vs SDC

SDC: $\hat{d}(x, y) = d(q(x), q(y)) = \sqrt{\sum_j d(q_j(x), q_j(y))^2}, \quad (12)$

ADC: $\tilde{d}(x, y) = d(x, q(y)) = \sqrt{\sum_j d(u_j(x), q_j(u_j(y)))^2}, \quad (13)$




symmetric case



asymmetric case

ADC vs SDC



| | SDC | ADC |
|--|----------------------------|---------|
| encoding x | $k^* D$ | 0 |
| compute $d(u_j(x), c_{j,i})$ | 0 | $k^* D$ |
| for $y \in \mathcal{Y}$, compute $\hat{d}(x, y)$ or $\tilde{d}(x, y)$ | $n m$ | $n m$ |
| find the k smallest distances | $n + k \log k \log \log n$ | |

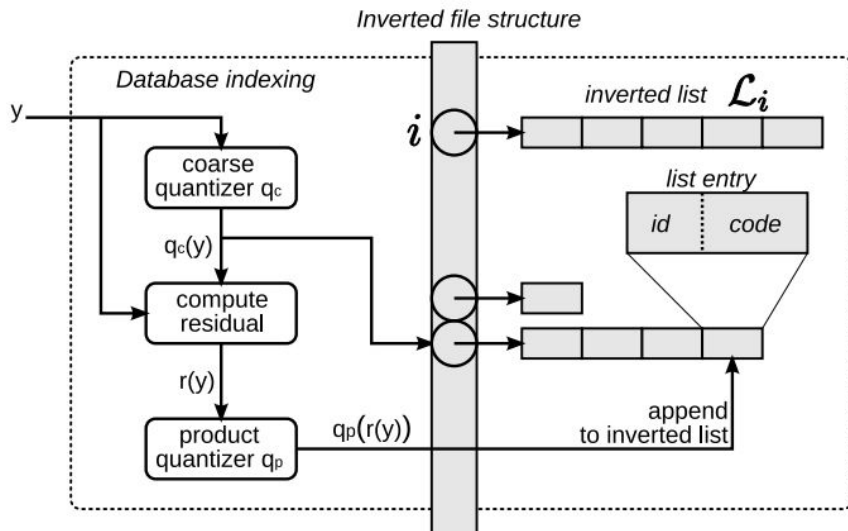
1. When n is large ($n > k^* D$), the most consuming operations are the summations in Equations 12 and 13.
2. The only advantage of SDC over ADC is to limit the memory usage associated with the queries, as the query vector is defined by a code.

Another problem: If n is large, an exhaustive search is prohibitive, as we need to index billions of descriptors and to perform multiple queries

IVFADC

Indexing a vector y proceeds as follows:

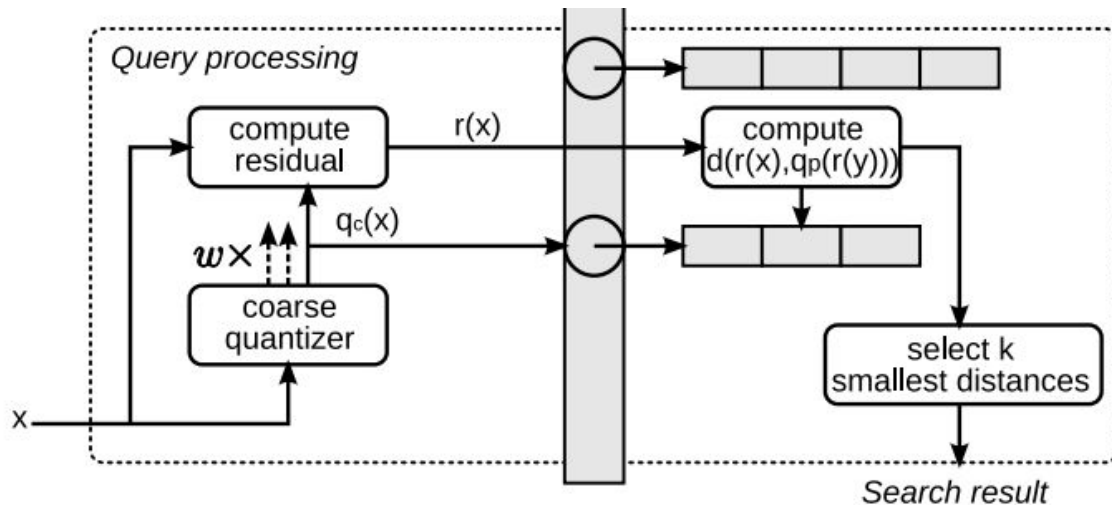
1. quantize y to $q_c(y)$
2. compute the residual $r(y) = y - q_c(y)$
3. quantize $r(y)$ to $q_p(r(y))$, which, for the product quantizer, amounts to assigning $u_j(y)$ to $q_j(u_j(y))$, for $j = 1 \dots m$.
4. add a new entry to the inverted list corresponding to $q_c(y)$. It contains the vector (or image) identifier and the binary code (the product quantizer's indexes).



IVFADC

Searching:

1. coarse quantizer
2. compute the residual $r(x)$
3. compute $d(x, y)$
4. select the k nearest neighbors



Experiments


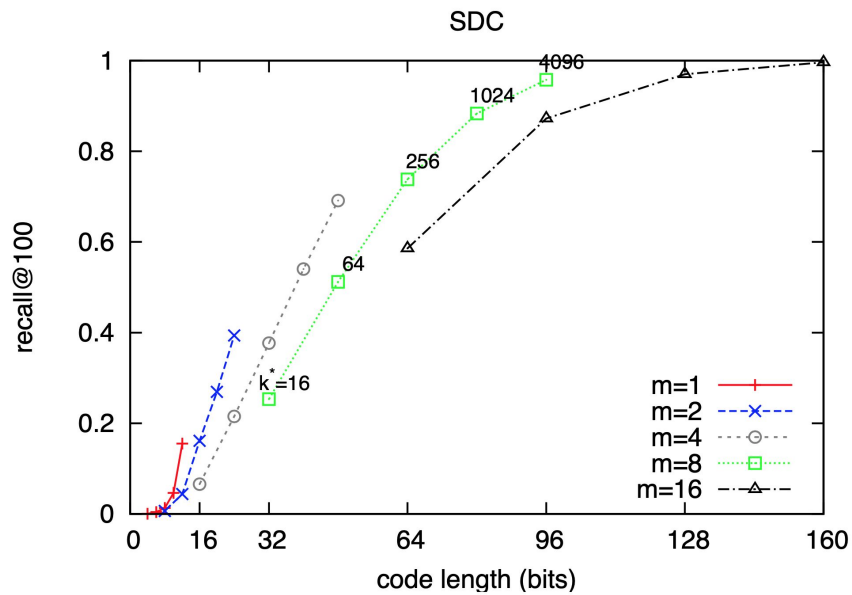
| vector dataset: | SIFT | GIST |
|-------------------------------|-------------|-------------|
| descriptor dimensionality d | 128 | 960 |
| learning set size | 100,000 | 100,000 |
| database set size | 1,000,000 | 1,000,991 |
| queries set size | 10,000 | 500 |

TABLE III


SUMMARY OF THE SIFT AND GIST DATASETS.

Experiments

SDC vs ADC





Experiments



Experiments

Compare with state of art method:


spectral hashing and **hamming embedding**



Experiments

Compare with state of art method:

FLANN



Experiments

Complexity and speed

| method | parameters | search time (ms) | average number of code comparisons | recall@100 |
|--------|-------------------|---------------------|---------------------------------------|------------|
| SDC | | 16.8 | 1 000 991 | 0.446 |
| ADC | | 17.2 | 1 000 991 | 0.652 |
| IVFADC | $k'=1\,024, w=1$ | 1.5 | 1 947 | 0.308 |
| | $k'=1\,024, w=8$ | 8.8 | 27 818 | 0.682 |
| | $k'=1\,024, w=64$ | 65.9 | 101 158 | 0.744 |
| | $k'=8\,192, w=1$ | 3.8 | 361 | 0.240 |
| | $k'=8\,192, w=8$ | 10.2 | 2 709 | 0.516 |
| | $k'=8\,192, w=64$ | 65.3 | 19 101 | 0.610 |
| SH | | 22.7 | 1 000 991 | 0.132 |

Experiments

large scale




Fig. 11. Search times for SIFT descriptors in datasets of increasing sizes, with two search methods. Both use the same 20 000-word codebook, $w = 1$, and 64-bit signatures.

Experiments

Image search




Fig. 12. Comparison of IVFADC and the Hamming Embedding method of [20]. mAP for the Holidays dataset as function of the number of distractor images (up to 1 million).

Thanks